

Uplift Note 2026-09-05 — 25 proposals from the Epistemic Fusion papers to glosa: what was wrong, what was fixed, what stands

Tier of this whole note: Dr (design synthesis / editorial reconciliation of prior reviews, not itself a new independent check). No Th_coqc claim is made anywhere below. No novelty/first/prior-art language. Comparisons are same/different/cited only.

สรุปภาษาไทย (10 บรรทัด)

- เอกสารนี้รวบรวมข้อเสนอทั้ง 25 ข้อจากการประชุม 2026-09-05 ที่นำแนวคิดจากเปเปอร์ Epistemic Fusion (v7.1/v8.1) เข้ามาสู่ glosa — ไม่ตัดทิ้งข้อไหน แม้จะถูกปฏิเสธในรอบแรก
- รอบแรก (MEETING_2026-09-05_judged.json) เก็บไว้เพียง 4/25 ข้อ เพราะผู้เสนอ 21 ข้อที่เหลืออ้าง id ป्लอม "C1-C11" ราวกับเป็น claim card จริง — grep แล้วไม่พบ id แบบนี้ในไฟล์ต้นฉบับเลย
- ไฟล์ต้นฉบับจริงเลข F1-F20 (v7.1) และหัวข้อ "Problem n"/"Repair n" (v8.1) เท่านั้น — "C1-C8" กลับเป็น namespace ที่ glosa เองใช้อยู่แล้วสำหรับ chair ruling คนละเรื่อง เกิดการชนกันโดยตรง ผิดกฎ 17
- คำสั่ง founder (BBL-2026-09-05-119): ห้ามลดค่ากล่าวอ้างให้ผ่าน ต้อง "ยกงานขึ้น" แทน — แก้เลขบรรทัด อ้างอิงจริงเปิดเผยช่องว่างที่แท้จริง ไม่ปิดเงียบ
- รอบสอง (SESSION_ARCH_v0.4_rejudged.json + _SPEC.md) ตรวจใหม่ทุกข้อด้วย citation แบบ file:line จริง ได้ผล: build_now 13 · with_revision 7 · still_open 4 จาก 24 ข้อ (ข้อ SA-5 หายไปจากการส่งต่อที่ตัดตอน)
- SA-5 (K_like state-machine citation) ถูกสืบหาใหม่ในโน้ตนี้เอง — grep พบตำแหน่งจริงใน EPISTEMIC_FUSION v7.1/v8.1 ครบ แต่เป้าหมายเดิม (§7.3) ผิดหัวข้อ ต้องแก้เป็น §4.4 เท่านั้น
- SA-5 ยังไม่ผ่านการตรวจสอบรอบอิสระเหมือนอีก 24 ข้อ (ตรวจโดย AI ผู้ร่างเอกสารนี้เอง = maker ตรวจ

- งานตัวเอง ไม่นับเป็น checker ตามกฎข้อ 3 ของ glosa) — จึงจัดสถานะ with_revision ไม่ใช่ build_now
8. รวมทั้งหมด 25 ข้อ: build_now 13 · with_revision 8 (รวม SA-5) · still_open 4
 9. พบ namespace collision จริง 2 จุด: (ก) "C1–C8" ของ chair ruling ชนกับ id ปลอมที่ผู้เสนอใช้ (ข) เลข kernel rule 29 ใน kernel/glosa_kernel.py ชนกับเลข flag-rule 29 ใน FOUNDATION_v0.6.md §3.3 คนละไฟล์คนละความหมาย ต้องเขียนแยกชื่อเต็มเสมอ
 10. เอกสารนี้เองมีข้อจำกัด: ไม่ได้ re-verify citation ทั้ง 24 ข้อจากไฟล์ต้นฉบับซ้ำอีกรอบ (อาศัยการตรวจของรอบสอง + สุ่มตรวจเอง) และยังไม่ผ่าน adversarial review อีกระลอกแน่นอนบน Zenodo — ต้องผ่าน gate นั้นก่อนเผยแพร่จริง
-

Status

- **K-state:** K0 (private candidate — this note has not been reviewed by anyone but the AI that drafted it and one direct-read spot-check pass by the same session; it is not K1 until a human timestamps/publishes it, and not socially stress-tested until an independent reviewer runs a second pass per glosa rule 3).
- **Tier:** Dr throughout, with per-row tiers as inherited from SESSION_ARCH_v0.4_rejudged.json (Dr/finite_diagnostic/Open, never Th_coqc) — see §2 table.
- **Novelty:** none claimed. This note re-anchors and reconciles existing review artifacts; it introduces no new mechanism except the SA-5 re-anchoring in §2, which is itself only a citation correction to an existing named target (design/FOUNDATION_v0.6.md §4.4).
- **Human owner:** Yaoharee Lahtee. The problem, the hypothesis-selection, and every founder decision named in §2/§7 of SESSION_ARCH_v0.4_SPEC.md remain human-owned; nothing here substitutes for that.
- **AI role:** drafted as assistant only. This note is a compilation/reconciliation pass, not a certification of any row's readiness — "build_now" below is the *reconciled reviewer verdict* carried over from SESSION_ARCH_v0.4_rejudged.json, not a

release decision. Per glosa rule 8, no independent check $\Rightarrow$ no release: none of the 25 rows are released by this note; each still needs its own DECISIONS.yaml row and, where flagged, a founder ruling before build.

- **Attribution:** no vendor/model name anywhere in this note or intended for the Zenodo record, per glosa rule 9 and CLAUDE.md's binding instruction; AI involvement is disclosed by role only.

1. What happened

Round 1 (design/MEETING_2026-09-05_judged.json, 25 rows) kept only 4 proposals (lrs.dr-tag-not-dag-node 8.0, HU-3 6.5, SA-2 6.0, SA-4 5.0) and refuted the other 21. The dominant refutation reason across those 21 was **citation fabrication**: proposers cited claim-card ids such as "C1", "C2", "C4", "C5", "C6", "C7", "C8", "C9", "C10" as if these were established evidence-card identifiers inside sources/notes/EPISTEMIC_FUSION_v7.1.txt / EPISTEMIC_FUSION_v8.1.txt. A direct grep run this pass confirms, again, that no such scheme exists in either source file:

```
grep -n "^C[0-9]" sources/notes/EPISTEMIC_FUSION_v7.1.txt
sources/notes/EPISTEMIC_FUSION_v8.1.txt
→ zero hits in either file
```

The source papers number findings F1–F20 (v7.1) or leave formulas/problems unlabeled under prose headings like "Problem 1", "Repair 8" (v8.1) — never C-anything. Worse, "C1"–"C8" is **already a live glosa id namespace**: design/CHAIR_RULING_v1.md:23-30 uses C1..C8 for chair rulings on independence_class, the disclaimer catalogue, fetch_status, verdict_tier, review_mode, non-collapse rows, lineage table, and origin_blackbox_ref — none of which is what any proposer meant. Citing "C4"/"C5"/"C7"/"C8" for session-architecture content therefore collided with an existing, unrelated meaning under Rule 17 (file:line + continuous verbatim citation).

Founder ruling BBL-2026-09-05-119 ("never lower a claim to make it pass — raise the work") set the correction discipline: every fabricated or off-by-N citation gets replaced by

its real file:line location, every genuinely weak mechanism gets its acceptance test tightened or its gap disclosed honestly, and no row is quietly dropped to shrink the problem.

Round 2 (design/SESSION_ARCH_v0.4_rejudged.json, 24 rows + design/SESSION_ARCH_v0.4_SPEC.md) re-anchored every citation to file:line with continuous verbatim text and re-ran the same citations_verified && control_mechanical && invariant_ok && score>=6 → build_now rule (one fixable defect → with_revision; a scope/mechanism conflict → still_open). Result: **13 build_now, 7 with_revision, 4 still_open** out of 24. **One proposal, SA-5, is absent from that 24-row file** — per CLAUDE.md rule 10 (added this same day, Blackbox Log BBL-2026-09-05-121, "No finding is ever dropped... after one proposal (SA-5) fell out of a re-anchoring pass through a truncated hand-off"), this is the exact incident that rule exists to name. §2 below re-anchors SA-5 directly against the source files, so this note's own table restores the full count of 25 and does not repeat the drop.

2. All 25 proposals

Columns: **claim as stated** is carried verbatim/near-verbatim from the proposal, unchanged by this note. **What was wrong (r1)** names the round-1 defect. **Correction** gives the real file:line and a verbatim excerpt (≤25 words) that replaces the fabricated citation. **Status now** is the round-2 reconciled verdict (build_now / with_revision / still_open), from SESSION_ARCH_v0.4_SPEC.md §3/§9, except SA-5 which this note re-anchors directly (marked below) and holds at with_revision pending an independent second check, per glosa rule 3.

id	lens	claim as stated
SA-1	session-architecture	Session has a name where AI-side mom

id	lens	claim as stated
		zero by scope-defin retained residue ma session's start state, should carry this.
schema.entry-resistance-precommit-field	kernel-gates	Precommitted resis (R*) at claim-card in block, declared befo candidate exists.
kernel.reciprocal-lineage-diagnostic	kernel-gates	Reciprocal-lineage c (χ_{recip}) over kg ed only, using session_id/actor/tim
schema.retention-direction-field	kernel-gates	Retention-direction (expansion/tunnel/u the hypothesis-selec persistence alone is validation.
SA-4	session-architecture	glosa explicitly decl decline, to add chi_ as schema fields no multi-turn evidence

id		lens	claim as stated
SA-3		session-architecture	exists and both diag defeat clauses.
			Human's problem s model, and verifica be captured before matching v8.1's Fog finding.
			Run the falsifier tes deferred (rule17.c certification-op mechanism, close th [Open] item.
SA-2	process.run-existing-open-falsifier-klike	kernel-gates	Retention/persisten habit across session evidence of expansi only an independen reclassify it. Default retained_direct
		kernel-gates	Session-boundary A reset assertion, meo checkable.

id	lens	claim as stated
schema.blackbox-question-trace	blackbox-dialogue	Blackbox Note records only twice (single-sl records how the liv changed turn-to-turn
schema.blackbox-line-origin-enum	blackbox-dialogue	Bare boolean ai_pr represent the name hazard — K_like s to K_assumed.
kernel.candidate-set-delta-cooking-step	blackbox-dialogue	Cooking log records but not the candida growth/collapse cou discipline demands
schema.blackbox-chooser-reaffirmation	blackbox-dialogue	became[] link reco candidate became s whether the human delegable selection

id	lens	claim as stated
schema.blackbox-language-bridge-subfield		reaffirmed it at reop
		translation-step c
	blackbox-dialogue	conflate a bilingual with a world↔read translation glosa's c distinguishes.
HU-1	human-uplift	Write the missing H Gate protocol card - points at nothing.
HU-2	human-uplift	Add an entry-anch to the intake stage, exposure.
HU-3	human-uplift	Session-close retent Blackbox Note cook states direction, not what was retained.

id	lens	claim as stated
HU-4	human-uplift	Advisor packet gain quality" self-check s recommending con
HU-5	human-uplift	P01_standpoint's "A completeness" gets bounded question s open-ended AI judg
lrs.dr-tag-not-dag-node	lrs-defeaters	Do not create a new for the LRS-defeater [Open] TODO line i P09 vetting.

id	lens	claim as stated
lrs.defeater-defeated-status-field	lrs-defeaters	Track defeater-attenuation. If a claimant has survived vs. never-tested, claim_card status, not tested.falsified
lrs.dialogue-table-claim-type-column	lrs-defeaters	Add a claim_type legitimate_defeat the LRS dialogue_
lrs.defeater-not-collapse-rule	lrs-defeaters	Kernel rule: reject "claim" or "feels solid" defeater, regardless
lrs.claim-type-defeater-enum	lrs-defeaters	Extend claim_type type-scoped defeater matching kernel rule
SA-5 (re-anchored this note, missing from	session-architecture	K_like state-machin

round 1 reused this exact letter-number pattern ("C4", "C5", "C7", "C8", "C9", "C10", "C1") for content from the source papers, which do not use that scheme at all (F1–F20 in v7.1; unlabeled "Problem n"/"Repair n" headings in v8.1). Avoidance: every corrected citation in §2 now uses `file:line`, never a bare letter-number id; where a source concept genuinely deserves a short handle for repeated reference (e.g. `chi_recip`, `retained_direction`), the handle is the concept's own name from the source text, not a re-minted Cn.

2. **rule29 collision.** `kernel/glosa_kernel.py`'s python-rule sequence stops at rule28 (confirmed by direct read this pass — no rule29/rule30 exist yet there). `FOUNDATION_v0.6.md` §3.3's own prose/disclaimer-style "flag-rule" sequence separately runs rules 26-28. `lrs.defeater-not-collapse-rule` and `lrs.claim-type-defeater-enum` (§2, both `build_now`) claim the next free `kernel/glosa_kernel.py` slots, rule29/rule30. `kernel.candidate-set-delta-cooking-step` (§2, `with_revision`) separately asks to extend `FOUNDATION_v0.6.md` §3.3's own sequence to "rule 29" — a same-numbered but non-colliding sibling in a different artifact/namespace. Avoidance: every reference to either sequence in this note and in any future `DECISIONS.yaml` row must write the full qualified name — `kernel.glosa_kernel.rule29` vs. `FOUNDATION_v0.6.md §3.3 flag-rule29` — never the bare number alone.

4. What the Epistemic Fusion papers gain from this pass

Now has a mechanical control in `glosa` (13 `build_now` rows, §2): the session-boundary AI-reset asymmetry (F8/F10/F12, v7.1) → SA-1's Blackbox Note fields; the pre-commitment of a checking route (R*, v7.1:351/356) → `schema.entry-resistance-precommit-field`; the direction- neutrality of retention (Gs/Ts/Δs, Repair 8, v8.1:341-361) → `schema.retention-direction-field` and its NC-77 non-collapse row; the full $H0=(P0,M0,U0,E0,\Phi0)$ entry-anchor apparatus (v7.1:344) → HU-2; the missing Human Mastery Gate wiring (already-real 10-question checklist, `templates/paper/arxiv-`

twocol/main.tex:375-378) → HU-1's new R8 publish-gate dimension; the bounded-completeness-prompt discipline (extending NC-58) → HU-5; and v8.1 §1's five-way claim-type/legitimate-defeater taxonomy (phenomenological/constitutive/structural_formal/diagnostic/empirical, v8.1:30-53) → the lrs-defeaters lens's five build_now rows (dialogue_table columns, defeater_log typing, kernel rules 29/30, the TODO-line filing discipline).

Stays Open, honestly: the reciprocal-lineage diagnostic chi_recip (v7.1:447-554) is build_now for its plumbing (a session_id field + a CLI subcommand) but its tier stays Open until a reproducibility fixture actually runs — it is never treated as a warrant, truth, or release gate (v7.1:496), matching the source's own framing of it as diagnostic-only. The momentum scores m^H/m^AI (v7.1:550-554) are explicitly **declined** as schema fields (SA-4), per their own documented self-defeat clause and glosa's current lack of any multi-turn/dialogue-history evidence mechanism — the papers' own hedge on these two constructs is upheld, not overridden. SA-2's hypothesis_selection.yaml application of the same Repair-8 direction-neutrality principle stays still_open until SA-1 supplies a session_id grouping key — the mechanism is sound but not yet executable on real multi-session data, which per BBL-2026-09-05-119 must be disclosed as an honest gap, not narrowed into a vacuous single-session-only test and called done. HU-3's session-close retention note (same Repair-8 principle, applied to the Blackbox Note's own cooking log instead of the hypothesis ledger) stays still_open for the same reason plus its own uncorrected citation defect. SA-3's pre-exposure H0-lite field and schema.blackbox-language-bridge-subfield both stay un-adopted in their current form — the former needs a single named schema target, the latter is withdrawn per a standing chair ruling its own proposer missed.

5. Defeaters for this note itself

1. **This note is a same-session compilation, not an independent check.** Per glosa rule 3 ("maker ≠ checker ≠ approver... same-model self-review has no standing"),

the reconciliation in §2 for the 24 rows carried over from `SESSION_ARCH_v0.4_rejudged.json` relies on that file's own two-skeptic-verified `citations_verified/control_mechanical/invariant_ok` fields rather than a fresh from-scratch re-grep of all ~90 citations across 24 rows by this session. This session did spot-verify a subset directly this pass (the `C1-C[0-9]` absence grep in §1; the `K_like` passages, `FOUNDATION_v0.6.md` §4.4/§7.3, and NC-47/NC-59 locations for SA-5 in §2) — those specific claims carry direct-read confirmation from this session; the remaining citations in §2 are relayed from round 2's own verification, not re-verified here, and should be read as such (a Dr-tier compilation of `finite_diagnostic-tier` upstream checks, not a fresh `finite_diagnostic` pass in its own right).

2. **SA-5's re-anchoring in §2 has had no independent second-skeptic pass**, unlike the other 24 rows. It is marked with `_revision` rather than `build_now` for exactly this reason, not because a defect was found in the citations themselves. A legitimate defeater to SA-5's with `_revision` status would be a second, independent read confirming the same file:line citations and the same §4.4-not-§7.3 target correction — until that happens, this row should not be treated as more settled than a solo pass warrants.
3. **This note has not passed the publish gate**. Per glosa rule 8 ("no independent check ⇒ no release... public push only through the publish gate (PR + adversarial review + leak scan)") and PUB-ADVERSARIAL-REVIEW's privacy/security-leak-scan requirement, this note must not be attached to the public Zenodo record (concept DOI 10.5281/zenodo.22318039; v7.1 = 22318040, v8.1 = 22319715) before an independent adversarial pass runs R1-R7 on it — including a leak scan for local paths/usernames and a check that no vendor/model name was introduced anywhere in this document, contrary to `CLAUDE.md`'s binding no-AI-attribution rule.
4. **The 13/8/4 count itself could be stale** the moment any of the 8 with `_revision` or 4 `still_open` rows is actually revised and re-scored — this note is a snapshot of the 2026-09-05 reconciliation state, not a live query against `DECISIONS.yaml`. A

future reader should re-run `python tools/anse_sync/cli.py list`
decisions (from the `internal command-center/`) rather than treat this table
as current beyond the date in its filename.

5. **No claim in this note has been through a founder ruling.** Every row marked
`founder_decision_needed: true` in the underlying
`SESSION_ARCH_v0.4_rejudged.json/_SPEC.md` (most `build_now` rows included
— see `SESSION_ARCH_v0.4_SPEC.md §7`) still requires a human decision before
any schema/kernel change lands; this note recommends nothing be built solely on
its own authority.

Rendered with `glosa render.py 0.1.0` — tier of content unchanged (Dr unless stated)