

Written by AI. Still True.

Knower Fetishism, Epistemic Pedigree, and the Human Face as a Bad Theory of Truth

Yaoharee Lahtee ¹

¹Open Civil Science Initiative, Thailand

August 2026

Abstract

Suppose a proposition survives source checking, inferential scrutiny, and independent reconstruction. You accept it. Then you learn that an artificial-intelligence system generated the draft, and your confidence falls. What, exactly, became epistemically worse? The familiar answer is that provenance can carry higher-order evidence about reliability. Correct. This paper grants that point and then asks what remains after provenance has done all the legitimate epistemic work available to it. The resulting problem is larger than AI. I argue that modern epistemology often slides from a plausible claim about the *possession* of knowledge—that knowing is a state or standing of a subject—to a much stronger and usually undefended claim about the *constitution* of epistemic value—that epistemically significant content, processes, records, or sources acquire standing only through a human or human-like knower. I call this the *Possession–Constitution Collapse*. Its social form is *knower fetishism*: the conversion of an earned relation among evidence, reliability, accountability, and inquiry into an apparently intrinsic property of the recognized knower. The paper develops a pre-subjective notion of *constraint structure*, stages the argument against knowledge-first epistemology, testimony, reliabilism, higher-order evidence, objective knowledge, information-theoretic epistemology, distributed cognition, and social epistemology, and proposes two constraints: provenance matters only through an identifiable epistemic channel, and source metadata must never be identified with the epistemic role for which it is merely a proxy. The thesis is not that machines know. It is that before anyone knew a fact there was already something to get right or wrong about it; the world did not wait for a human knower before acquiring the differences, traces, dependencies, and constraints to which later knowledge must answer. A non-knower can therefore occupy an epistemically productive role without being promoted into personhood. Institutions remain indispensable, but their authority is derivative from epistemic work, not magical pedigree. AI is philosophically useful because it removes the human face from a claim and reveals how much epistemic credit we had quietly attached to the face. The paper’s provocation is deliberately asymmetric: if one wishes to reserve *knowledge* for a subject’s successful possession, one may keep the word; one may not infer from that linguistic reservation that the subject created the truth, trace, dependence, constraint, or record to which successful possession must answer.

Keywords: epistemology; artificial intelligence; objective knowledge; testimony; epistemic authority; anthropocentrism; provenance; scientific instruments; social epistemology

Disclosure challenge. Generative AI assisted with literature mapping, adversarial argument testing, language editing, and LaTeX drafting for this manuscript. The author selected, checked, revised, and endorses the final argument. If this disclosure lowers your assessment, do not merely report the discomfort. Identify the proposition, inference, evidential relation, or reliability condition that became worse when you learned it.

1. The sentence was fine five seconds ago

Take a deliberately clean case. A reader, *S*, receives a proposition *p*, an argument *A*, and an evidential record *E*. She checks the cited materials, reconstructs the relevant inference, and finds no defect. She judges *p* well supported. A moment later she learns only this:

The draft was generated by an AI system.

Suppose her confidence falls. The cheap response is that she has committed a genetic fallacy. That response is too cheap. Source information can be epistemically relevant. Reliabilism, expertise, testimony, higher-order evidence, and genealogical debunking all explain why learning something about the origin of a belief can rationally alter confidence without altering the proposition’s

truth-value [9, 11, 10, 2, 14]. If “AI-generated” reveals fabricated references, hidden dependence, systematic error, non-independent verification, or a previously unmodeled failure mode, confidence may properly fall.

So let provenance win its strongest case. The harder question begins afterward:

After provenance has told us everything epistemically relevant that provenance can tell us, what exactly is the label still doing?

The question matters because two different objects are easily collapsed: the epistemic properties of a claim and the social pedigree of the route by which the claim arrived. A useful proxy for reliability can become so familiar that we begin treating the proxy as though it

were reliability itself.

The joke is simple. Nothing happened to the sentence. The punctuation survived. The evidence survived. The inferential relations survived. Only the biography changed.

If that biography changes the verdict, it may do so for excellent reasons. But those reasons must be named. A human face is not a free epistemic variable.

But that joke points toward a larger embarrassment. Epistemology has spent centuries asking what a knower must be like, what a knower may believe, when a knower is justified, whom a knower may trust, and how a knower acquires knowledge. Those are legitimate questions. They become less innocent when the grammar of *someone knows that p* quietly becomes an ontology in which the knower is treated as the source of epistemic standing itself.

This paper asks whether we have confused the conditions under which knowledge is *possessed* with the conditions under which reality is *epistemically structured*.

2. Before anyone knew it, there was already something to get wrong

Before there were physicists, stars behaved in ways that later made some astrophysical claims correct and others false. Before there were paleontologists, animals died, bones fractured, sediments accumulated, and traces were retained. Before there were geneticists, inheritance displayed stable dependencies. None of these events filed an interpretation with a human observer before occurring.

The claim is intentionally weaker than “knowledge existed before knowers.” That slogan invites a terminological fight and wins nothing we need. If “knowledge” is reserved for a factive state or standing possessed by an epistemic subject, then let the subject keep the noun. Williamson’s knowledge-first program gives that subject-involving state an especially fundamental role [21]. The present argument can grant the usage completely.

What it refuses is the inflationary inference from *someone must possess knowledge to nothing can be epistemically structured before someone possesses it*. Before anyone knew that *p*, there could already be a fact of the matter that made *p* correct rather than $\neg p$, traces selectively dependent on that fact, and constraints capable of defeating a later interpretation.

Before anyone knew it, there was already something to get wrong.

That sentence is the realist minimum of the paper. A theory may reject it, but then the dispute is no longer primarily about artificial intelligence. It is about whether inquiry discovers constraints that precede inquiry or partly constitutes the very correctness conditions by which inquiry is judged.

2.1. The Jurassic test

Imagine a dinosaur whose femur fractured long before any human existed. The fracture generated a physical

pattern. Millions of years later, investigators disagree about whether the injury occurred before or after death. One interpretation fits the retained structure better than the other.

Now ask four questions.

1. Did the fracture occur, if it occurred, before paleontology?
2. Did the resulting structure differ depending on what happened?
3. Can that retained difference make a later paleontological claim better or worse?
4. Did the difference wait for a human interpretation before being the difference it was?

A “yes, yes, yes, no” answer is enough. The fossil need not know anything. It need not carry a proposition in English. It need only be the sort of worldly remainder against which later representations can succeed or fail.

A fossil does not need a theory of knowledge to make your theory of the fossil wrong.

Reality can say no without speaking.

Dretske’s information-theoretic epistemology is useful here because information, in his technical sense, is objective: a signal carries information in virtue of lawful dependence on a source condition, not because a human first interprets it [4]. Popper pushed in a different direction with the deliberately provocative project of an “epistemology without a knowing subject,” distinguishing subjective states from objective contents of thought and scientific theories [18]. Frege’s realism about thoughts likewise denies that truth-bearing thought contents are private mental ideas owned by an individual thinker [6]. Baird goes further in a material direction, arguing that scientific instruments can embody or constitute forms of “thing knowledge” rather than merely assisting propositional thought [1].

These views are not equivalent. Dretske’s information is not Popper’s World 3. Fregean thoughts are not spectrometers. Baird’s instruments do not prove that fossils literally know anything. Their disagreement is useful precisely because it blocks an easy escape: epistemic significance has never had only one philosophically uncontested address.

To avoid winning by vocabulary, introduce a deliberately weaker notion.

Definition 1 (Pre-subjective constraint structure). *A pre-subjective constraint structure is a worldly state, relation, trace, record, or lawful dependence capable of constraining which later representations succeed or fail, independently of whether any actual subject has yet represented or interpreted it. The adjective epistemic is deliberately withheld at this stage: the structure need not already be evidence-for-us, meaningful-to-us, or anyone’s knowledge in order to constrain what later inquiry can responsibly say.*

Interpretation matters for uptake, representation, comparison, inference, and use. It does not travel backward in time to manufacture the event, dependence, or

trace that constrains the interpretation. The point can therefore be stated without giving a fossil a mind:

If you insist that knowledge must be owned, keep the word.

The world does not care.

What makes ownership epistemically answerable was there before the owner.

3. The Possession–Constitution Collapse

Let a minimal **Subject Condition** be:

$$\text{SC} : K(S, p) \rightarrow \text{Subject}(S).$$

If S knows p , some subject S occupies the knowing relation. This is deliberately weak.

Now consider a stronger thesis:

$$\text{HSC} : \text{Epi}(X, p) \rightarrow \text{Knower}(X, p),$$

where $\text{Epi}(X, p)$ means that X plays an epistemically significant role in a route by which p becomes warranted, available, retained, or inferentially constrained.

Call this the **Knower–Constitution Condition**. In its human form, it becomes the Human-Source Condition: epistemically serious production must terminate in, originate from, or pass through something sufficiently like a recognized human knower.

Nothing in SC entails this stronger condition.

Definition 2 (Possession–Constitution Collapse). *The Possession–Constitution Collapse occurs when a condition on who can possess knowledge is treated as a condition on what can constitute, constrain, retain, transmit, or generate epistemically significant structure.*

The collapse is tempting because ordinary human cases package many roles into one body. A scientist may generate a claim, understand it, endorse it, answer for it, possess relevant expertise, and carry institutional standing. Because the bundle is historically common, philosophy can mistake co-location for identity.

Principle 1 (Role Separation). *Generation, truth, evidential support, reliability, understanding, possession, endorsement, accountability, credibility, and institutional authorization are distinct epistemic roles. No pair may be identified merely because one human agent often occupies both.*

The principle is not anti-human. It is anti-shortcut. A human being can occupy many roles at once. What she cannot do is make those roles identical by occupying them.

**The human body was a convenient interface.
Epistemology mistook the interface for the architecture.**

The collapse is easy to commit because ordinary grammar hides it. “Alice knows that the sample contains lead” naturally puts Alice at the center. The grammar says almost nothing about where the epistemically load-bearing work occurred: in the calibration procedure, the detector, the sample preparation, the chemical regularity, the laboratory archive, the inferential model, the testimony of technicians, or Alice’s final endorsement.

Giere’s analysis of distributed cognition in epistemic cultures makes this division explicit: scientific cognition can be distributed across people and artifacts even if one reserves epistemic agency for human components [8]. Humphreys makes the anti-anthropocentric point even more sharply for computational science: human cognitive capacities need not remain the ultimate standard of epistemic success [12, 13].

Proposition 1 (No transitivity of knowerhood). *From the fact that knowledge has a subject it does not follow that every epistemically significant node in the route by which that subject comes to know must itself be a knower.*

A thermometer need not know the temperature. A telescope need not believe in galaxies. A calculator need not understand multiplication. A simulation need not possess the theory it helps investigators evaluate. The absence of knowerhood at an intermediate node is therefore not, by itself, an epistemic defect.

The burden moves. If a route is defective, find the defect in evidence, reliability, calibration, dependence, opacity, selection, inference, accountability, or some other relation that bears epistemic weight.

“But it does not know” is not yet an argument.

3.1. The source label is a readout, not an oracle

The argument can be made more precise without pretending that finite inquiry has transparent access to reality. Let a worldly condition be W , let an inquiry procedure or access operator be $\mathcal{O}_A$, and let R_A be the record available to agent or system A under a selection policy Π_A :

$$R_A = \mathcal{O}_A(W; \Pi_A).$$

The notation is intentionally modest. It says that what an investigator receives is an operator-conditioned record of the world. It does *not* license the identification $R_A = W$. Human observers do not escape this merely by having biographies.

Now let $m(A)$ be metadata about the source—“senior professor”, “top journal”, “AI-generated”, “anonymous amateur”—and let $\hat{\rho}(A)$ be the reliability estimate an evaluator extracts from that metadata. Metadata can be excellent evidence about reliability. But:

$$m(A) \neq \rho(A),$$

and neither metadata nor estimated reliability is identical to the epistemic standing of p .

The mistake targeted in this paper is therefore a *role-identification error*: a proxy that is useful in one inferential role is promoted into the thing it was introduced to estimate.

Principle 2 (Bridge Burden). *Any inference from source metadata to a change in epistemic standing must identify the mediating relation that does the work—for example reliability, independence, access, assurance, accountability, calibration, or error correlation. Where no such bridge is supplied, the inference is pedigree substitution rather than source-sensitive epistemology.*

Table 1: Ten roles routinely compressed into the figure of “the knower.” The table is diagnostic: co-location in one person does not establish identity among the roles.

Role	Governing question	What it is not
Truth	Is p the case?	Not conferred by authorship, prestige, or recognition.
Constraint	What worldly difference can make a later representation fail?	Not yet evidence-for-us or anyone’s possession.
Evidence	What bears on whether p should be accepted?	Not identical to a source label.
Reliability	How well does a process discriminate success from error across cases?	Not the truth-value of this particular output.
Provenance	By what route did this representation arrive?	Not a verdict on the representation by itself.
Credibility	How much epistemic weight is assigned to a source?	Not automatically deserved authority.
Accountability	Who can answer, correct, retract, or defend?	Not a truth-maker and not identical to generation.
Legitimacy	What social or institutional standing is granted?	Not the same as evidential support.
Authorship	Who is properly answerable for the scholarly act?	Not necessarily who generated every token.
Possession	Who, if anyone, knows p ?	Not what constituted every epistemically productive step on the route.

Principle 3 (No Bare Pedigree). *A source label is epistemically incomplete reporting. The relevant object is not merely who or what produced the claim, but the tuple of production procedure, selection conditions, dependencies, checks, error model, inferential role, and answerability that makes the source cue evidentially useful.*

This is the source analogue of a familiar scientific discipline: never let the displayed value erase the measurement procedure that gave it meaning. “Professor” and “AI” are not epistemic magnitudes with values printed on their foreheads. They are compressed descriptions of routes. Compression is allowed. Transubstantiation is not.

A PhD is not a truth-maker.
An AI label is not a defeater by sacrament.
Show the channel.

4. Knowledge-first meets a world that came first

Knowledge-first epistemology gives us the strongest possible subject-centered opponent. Williamson treats knowledge as a fundamental factive mental state rather than as an analysable composite of belief, truth, and justification [21]. The present argument need not deny that knowing is a state of the subject. It challenges a different inference: that because *knowledge-as-state* is fundamental in explaining belief, evidence, assertion, or action, the subject who occupies that state thereby becomes constitutively prior to all epistemically significant structure.

Knowledge-first may put knowledge first in an explanatory theory of mind. It does not follow that the knower came first in the history of the world, or that human-like sourcehood comes first in the architecture of inquiry. The order of explanation is not the order of arrival. A dinosaur is not obligated to wait for epistemology before breaking its leg, and a detector is not obligated

to acquire beliefs before its output can embarrass a theorist.

That inference would be surprising. Williamson himself emphasizes that what can be known does not exhaust what is true. The world therefore exceeds the set of actual epistemic possessions. The factive character of knowledge already points outward: knowing succeeds because the subject is appropriately related to how things are, not because the subject manufactures how things are.

Popper’s objective epistemology moves almost in the opposite direction. Scientific knowledge, on his view, can be considered at the level of objective contents, problems, arguments, and theories rather than reduced to psychological states of knowers [18]. Popper’s World 3 is historically produced by minds, so it does not establish the stronger claim that knowledge literally predates every knower. What it does establish is enough for the present dialectic: epistemology can coherently investigate epistemic contents and relations without making current possession by a subject constitutive of their objective status.

Dretske then removes another layer of dependence. Information can be objective before cognition takes it up; knowledge is an information-dependent state, not the creator of the informational relation [4]. Baird removes another: material artifacts can embody epistemic achievements that exceed what any individual presently has in mind [1]. Giere and Humphreys remove another: epistemic production is heterogeneous and technologically extended [8, 12].

The resulting picture is not that subjects disappear. It is that the subject arrives *inside* an already structured problem.

The knower may be necessary for knowing.
The knower is not therefore the author of what makes
knowing correct.

This distinction is the paper’s deepest anti-

anthropocentric commitment. We once removed the human from the center of the cosmos and from a special biological creation story. Epistemology may still preserve a smaller throne: reality can be independent of us, but epistemic legitimacy somehow requires our face.

5. Source matters. Unfortunately.

Any attempt to decenter the knower becomes unserious if it ignores why sources matter. Hardwig’s epistemic dependence is unavoidable in modern inquiry: no individual can personally reproduce the chains of expertise on which scientific belief depends [11]. Goldman shows why novices rationally use evidence about experts, track records, agreement, and competence when direct access to the subject matter is limited [10]. Coady and Lackey establish the depth of testimony as a source of knowledge [3, 16]. Zagzebski gives a sophisticated account of epistemic authority rather than treating deference as mere irrational submission [22].

Higher-order evidence makes the point sharper. Christensen shows that information about one’s epistemic performance can rationally alter confidence even when first-order evidence is unchanged [2]. Kahane’s analysis of evolutionary debunking arguments shows how genealogical facts can undermine justification without changing the truth-value of the target proposition [14].

So the paper concedes all of the following:

- source identity can be evidence;
- credentials can rationally compress track-record information;
- human testimony can carry assurance and responsibility unavailable from an impersonal process;
- institutional pedigree can indicate methods, training, correction, and accountability;
- AI provenance can rationally lower confidence when it identifies a real failure mode.

Good. Let every opponent win those points. The paper needs none of them to lose.

Now the question can finally become difficult.

If the epistemic force of pedigree is explained by reliability, evidence, assurance, accountability, and institutional procedure, what epistemic work remains for pedigree *after those relations have done theirs*? If the answer is “humanity”, the burden is no longer rhetorical. Specify the channel by which humanity improves the epistemic relation in the case at hand. If no channel can be specified, the word is functioning as an honorific, not an explanation.

6. Knower fetishism

The term “fetishism” is intentionally accusatory, but the accusation is structural rather than psychological.

Definition 3 (Knower fetishism). *Knower fetishism is the treatment of the recognized identity, humanity, credentials, or social standing of a knower as though it possessed epistemic value in itself, after abstracting*

away the evidential and relational work through which that status earns epistemic relevance.

The structure resembles familiar cases of reification. A relation produced by training, successful performance, institutional checking, criticism, and accountability is compressed into a status term: *expert*. Compression is useful. The mistake begins when the status term is allowed to behave like an intrinsic epistemic property.

Professor *P* may deserve high prior credibility because *P* has an excellent track record, relevant expertise, transparent methods, and exposure to criticism. But if we subtract those relations and retain only “Professor *P*”, what remains epistemically? A title. If the title still moves credence after the reasons encoded by the title have been held fixed, the compression format has begun impersonating the data.

Likewise, “AI” may rationally trigger caution because a system has known failure modes. But once the relevant failure modes are represented, an unexplained residual penalty cannot be defended merely by repeating the species or substrate of the source.

There is a social reason the mistake is attractive. A recognizable human knower is an administratively convenient address for trust, blame, credit, responsibility, and legitimacy. That convenience is real. It is not yet epistemology. The danger begins when the address of responsibility is mistaken for the address of truth.

Perhaps we never wanted truth alone.

We wanted truth with a human witness.

The problem is not that we respected knowers too much.
The problem is that we forgot what the respect was for.

This is why AI is a useful philosophical irritant. The traditional human author bundled generation, belief, assertion, intention, responsibility, reputation, and social identity into one body. Generative systems unbundle those roles. Once the bundle comes apart, epistemology has to decide which role was actually doing the work.

7. The Epistemic Pedigree Paradox

Consider two versions of the same checked argument. Version *H* is labeled “written by a senior professor.” Version *M* is labeled “drafted by AI and independently checked by an accountable human.” Suppose the proposition, cited evidence, derivation, verification procedure, known dependencies, and endorsing responsibility are held fixed as far as the case allows.

If credence still differs merely with the label, we have a residual provenance effect. More generally, the epistemically relevant object is a route, not a badge:

(origin, procedure, dependence,
checks, answerability) $\rightarrow \Delta\text{Cr}(p)$.

A bare label suppresses the very variables that could justify its influence.

For two provenance descriptions O_1 and O_2 , a useful isolating diagnostic is

$$\text{RPE} = \text{Cr}(p \mid E, R, A, O_1) - \text{Cr}(p \mid E, R, A, O_2),$$

where E denotes first-order evidence, R the represented reliability and dependence conditions, and A answerability or responsible endorsement. The expression is not a psychological law. It asks what, normatively, remains for provenance to do once the channels invoked on its behalf have been made explicit.

Definition 4 (Residual Provenance Effect). *A Residual Provenance Effect occurs when epistemic assessment changes with provenance after the epistemically relevant pathways through which provenance could rationally matter have been held fixed or explicitly represented.*

The point is not that real cases will perfectly screen off every pathway. They usually will not. The point is normative: any residual difference incurs a burden of explanation.

Principle 4 (Provenance Relevance Constraint). *Provenance may rationally alter epistemic standing insofar as it changes total evidence, reliability, dependence, assurance, accountability, or another specified epistemically relevant condition. Provenance has no additional free-standing force merely by redescribing the origin of an otherwise fixed epistemic route.*

The formal ancestry is intentionally unglamorous: the structure resembles screening-off. Once the variables through which a common cause bears on an outcome are adequately conditioned on, the residual label should not acquire unexplained causal or evidential force [19]. The philosophical problem is therefore not inventing conditional independence. It is noticing when social categories are allowed to bypass it.

This principle is deliberately conservative. It does not say provenance decays monotonically with verification. A late discovery of data leakage, fraud, shared model dependence, or correlated error can make provenance newly relevant. It says only that relevance must travel through an epistemic channel.

The paradox appears because pedigree is most useful when direct checking is difficult. We invented expert shortcuts because we could not inspect everything. Yet after inspection becomes more possible, we can find ourselves inspecting the shortcut instead.

We needed credentials because we could not check everything.
Then we learned to check more things and began checking the credentials instead.

7.1. The author-shaped placebo

Consider two texts that are token-for-token identical. One is labeled “Professor.” The other is labeled “AI.” If the labels rationally signal different reliability, dependence, or accountability, different credence may be warranted. That is ordinary source-sensitive epistemology. But once those pathways are fixed, any remaining difference is not a mysterious extra property of the text. It is a residue in the evaluator.

This is why the paper’s provocation is conditional rather than populist. It does not demand equal trust for equal strings. It demands an account of the *difference-maker*. If the answer is reliability, show the reliability

difference. If it is independence, show the dependence. If it is responsibility, show what responsibility changes for the epistemic task. If it is merely “one came from a human,” then humanity has ceased to be a description and started functioning as an unexplained epistemic honorific. If “human” is a reliability variable, estimate the reliability. If it is a moral category, do not smuggle it into credence. Cheekbones are not calibration.

A name can be evidence. It cannot be evidence of itself forever.

8. AI as philosophical acid, not as a new person

Much contemporary debate asks whether conversational AI can testify, assert, understand, or count as an epistemic authority. Freiman argues that beliefs acquired from conversational AI do not fit neatly into the traditional instrument/testimony dichotomy [7]. Elizabeth Fricker argues forcefully that human and AI “testimony” differ metaphysically and epistemically, with responsibility playing a central role in genuine human testimony [5]. Sahebi and Formosa analyze the dilemma created by epistemic distrust of AI-mediated communication: caution can be rational while indiscriminate penalties can also distort human epistemic relations [20].

The present paper takes an impolite shortcut through that debate.

Suppose current AI does not know. Suppose it does not believe. Suppose it does not testify. Suppose it has no understanding worth the name. Suppose it has no moral standing. Grant the skeptical side everything.

The central argument survives.

Because the question is not:

Can the machine be promoted into the club of knowers?

It is:

Why must it be a knower in order for us to know through it?

The thermometer was never promoted. Neither was the telescope. Neither was the archive. Neither was the calculator. Epistemology already permits non-knowers to occupy epistemically productive roles. AI is difficult because it performs those roles in language, the medium we had culturally associated with minds, assertion, testimony, and authorship.

AI therefore acts as philosophical acid. It dissolves a bundle that used to arrive in one human body and leaves separate questions on the table: Who generated the string? What supports the claim? What process produced it? Who checked it? Who endorses it? Who can answer objections? Who bears responsibility? Which parts are reproducible? Which dependencies are hidden?

The human author once answered all of these questions by standing there.

That was convenient. It was never an argument that standing there was itself epistemically magic. The machine is philosophically embarrassing not because it

lacks a soul, but because it forces us to itemize what the soul was allegedly doing for the inference.

Written by AI. Still true.

The first sentence concerns genealogy. The second concerns the world.

A philosophy of knowledge should not erase the punctuation without an argument.

The paper’s smallest philosophical object may therefore be the period in its title. A full stop is doing the work that source-sensitive reasoning often forgets to do: keeping two questions apart until a warranted bridge connects them.

9. Credibility: when the face becomes evidence

Miranda Fricker’s account of testimonial injustice is indispensable because it shows that credibility is not distributed in a social vacuum. Speakers can receive credibility deficits because identity prejudice enters epistemic uptake [17]. The lesson for the present argument is structural, not analogical. An AI system is not thereby a victim of testimonial injustice in the morally relevant sense. The interesting point is that audiences already use social identity as part of epistemic assessment, sometimes legitimately and sometimes not.

The human face can be evidence. A known scientist’s face may stand in for a rich history of competence and accountability. A marginalized speaker’s face may, unjustly, trigger prejudice. The same visible cue can therefore carry epistemically relevant and epistemically corrupting effects through different mechanisms.

Credibility and legitimacy must nevertheless remain separate. Credibility concerns the weight actually or rationally assigned to a source; legitimacy concerns the standing a social or institutional order grants to that source or claim. A claim can be epistemically strong and institutionally illegitimate. A claim can be institutionally legitimate and false. A proposition may have a passport and still be wrong.

This is why the paper refuses the slogan “ignore the source.” The demand is harder:

For every source effect, name the epistemic relation that earns it.

Once that discipline is imposed, the authority of the knower becomes decomposable. What looked like an intrinsic property of a person becomes a network of relations: competence, track record, access, independence, honesty, answerability, methodological discipline, and exposure to correction.

Knower fetishism is what happens when we stop decomposing.

10. Institutions after epistemic monarchy

The attack on knower fetishism is not an attack on institutions. It is an attempt to give institutions a better reason to exist.

Hardwig’s epistemic dependence and Kitcher’s division of cognitive labor show why modern knowledge

cannot be organized as universal first-person verification [11, 15]. Universities, journals, laboratories, professional bodies, and archives coordinate expertise, preserve memory, train judgment, distribute criticism, and sustain correction across time. Their epistemic value can be enormous.

But enormous is not mysterious.

Principle 5 (Friction, not magic). *Institutional certification has epistemic force insofar as institutions produce reliable epistemic friction: criticism, validation, replication, robustness testing, archival continuity, correction, and answerability. Certification does not create truth by status alone.*

As claim production becomes cheaper, the institutional problem changes. Scarcity of production can no longer do as much filtering work. The appropriate response is not to make entry artificially scarce. It is to make survival under criticism expensive.

Institutions should be difficult places for claims to survive, not difficult places for claims to enter.

The alternative generates a regress. If evidence counts only after an authorized institution recognizes it, what epistemic fact authorizes the institution? If the answer is that its procedures reliably distinguish better from worse claims, then procedure does the work. If the answer is simply that the institution is authorized, epistemology has quietly become administration.

10.1. The Verification Paradox

Suppose a claim is denied standing until an institution certifies it. If certification adds criticism, replication, methodological checking, or accountability, there is no paradox: the institution added epistemic work. But if every epistemically relevant feature is already fixed and the remaining difference is only the stamp, then certification is recognizing permission rather than producing warrant. If the claim still cannot count without the stamp, “knowledge” has begun to track authorization.

The paradox does not show that certification is indispensable. It forces certification to disclose its difference-maker. The institution may be an indispensable machine for producing epistemic friction. It is not a sacrament.

Authority should be the receipt for epistemic work, not the merchandise. Peer review can expose a claim to disciplined criticism; it cannot transubstantiate a false proposition into a true one. A journal can reject a truth and publish an error. Reality keeps no impact-factor database.

11. Fourteen objections, none invited to be polite

11.1. Objection 1: “Facts are not knowledge”

Correct. Under ordinary subject-involving usage, a dinosaur fossil existing before inquiry is not someone’s knowledge.

Reply. That concession is built into the argument. The claim is that conditions of *possession* do not exhaust

conditions of *epistemic constraint*. Calling the pre-subjective structure “not knowledge” does not make the structure begin later. If the objection is merely lexical, it wins the noun and loses the ontology.

11.2. Objection 2: “Knowledge-first already says knowledge is world-involving”

Williamson’s factive mental state is not a Cartesian item sealed inside a head. Why accuse knowledge-first epistemology of anthropocentrism?

Reply. The paper does not accuse knowledge-first theory of idealism. It targets a slide in surrounding epistemic practice: from knowledge being a subject’s world-involving state to human-like subjecthood being treated as a qualification for epistemic contribution. The former does not entail the latter.

11.3. Objection 3: “Popper’s objective knowledge is still human-produced”

World 3 does not give you pre-human knowledge.

Reply. Agreed. Popper is used to establish that objective epistemic contents can be studied without current possession by a knowing subject. The pre-human claim is weaker and realist: facts, traces, lawful dependencies, and constraints precede our uptake of them. The paper does not smuggle World 3 into the Jurassic period.

11.4. Objection 4: “Information is not knowledge”

Dretskean information can exist without a knower, but knowledge requires the right information-dependent state.

Reply. Exactly. Information demonstrates that epistemically load-bearing structure need not originate in a knower. Knowledge can be the successful uptake of structure that did not wait for uptake to exist.

11.5. Objection 5: “AI is not a thermometer”

LLMs are generative, opaque, semantically flexible, and capable of fluent fabrication. Instrument analogies understate the danger.

Reply. Correct. The analogy proves only a negative: non-knower status alone cannot disqualify an epistemic contributor. AI-specific defects must be identified at the level of reliability, dependence, evidence, or accountability. That is stricter than anthropomorphic dismissal, not more permissive.

11.6. Objection 6: “Provenance is higher-order evidence”

Learning that AI generated a claim can rationally lower confidence even after first-order checking.

Reply. Yes. The Provenance Relevance Constraint preserves exactly that possibility. The target is unexplained residual force after the relevant higher-order pathway is represented. This objection kills a naive genetic-fallacy argument. It does not kill the present one.

11.7. Objection 7: “Unknown unknowns never screen off”

No finite evaluator can know that every relevant AI failure mode has been represented. Machine provenance may therefore rationally lower confidence because it signals risks we have not yet discovered. The residual-provenance test appears impossible to satisfy.

Reply. The objection limits empirical identification; it does not license a source label to become a universal solvent. If “AI-generated” rationally predicts unmodeled failure, then the operative reason is uncertainty about reliability. Say that. Estimate it where possible. Preserve it where it cannot be estimated. What cannot be allowed is an unspecified remainder that defeats any amount of contrary evidence merely because its content is unknown. Ignorance is allowed to be uncertainty. It is not allowed to cosplay as evidence. An unfalsifiable pedigree penalty is not caution; it is immunity from correction.

11.8. Objection 8: “Your independent checks may share the same ancestor”

A claim can appear independently verified while the generator, verifier, retrieval system, benchmark, or dataset share a hidden common source. Two checks may therefore be one epistemic route wearing two interfaces.

Reply. Correct. Independence is a property of the route, not the number of reviewers or buttons at its end. A pipeline that shares model family, training contamination, corrupted data, or a common inferential bottleneck does not screen off provenance. This strengthens the Bridge Burden: a verification claim must carry its dependency structure. Two checks are not two checks merely because they have two buttons.

11.9. Objection 9: “Accountability is irreducibly human”

Scholarly knowledge requires agents who can answer, retract, correct, and bear responsibility.

Reply. Accountability may remain human. Generation need not. Verification need not. Calculation need not. Retrieval need not. The point of role decomposition is precisely to preserve accountability without pretending that the accountable person personally performed every epistemically productive operation.

11.10. Objection 10: “This is just conditional independence”

Once every relevant pathway from provenance to reliability is screened off, of course the label should not matter. The thesis is formally trivial.

Reply. Good. Sometimes philosophy’s most useful job is to show that a socially sophisticated practice rests on a formally embarrassing confusion. Simplicity is not always a weakness. Sometimes it is the indictment.

11.11. Objection 11: “Without the knower, anything goes”

If human status loses its privilege, rumor, machine output, amateur speculation, and expert judgment collapse

into one epistemic level.

Reply. The opposite follows. Remove automatic human credit and explicit epistemic standards become more important: evidence, calibration, track record, independence, reproducibility, error detection, answerability, and criticism. The proposal is not equal credibility. It is equal liability to epistemic testing.

11.12. Objection 12: “There is no epistemic structure without interpretation”

On a sufficiently strong constructivism, traces become evidence only within concepts, practices, and interpretive schemes. Calling anything epistemic before those practices arise therefore risks smuggling human semantics into a pre-human world.

Reply. This is why the paper now calls the pre-subjective item a *constraint structure*, not evidence and not knowledge. The paper can concede that *evidence-for-us* requires uptake within a practice. It needs only the weaker claim that uptake is answerable to differences it did not create. The fracture pattern, stellar emission, or causal dependency need not already mean anything in our vocabulary. It need only constrain which later meanings survive contact with it. Interpretation supplies a mapping; it does not thereby manufacture every constraint on the mapping. If that weaker realism is rejected, say so. But then AI is no longer the controversial part of the paper.

11.13. Objection 13: “Your formalism just hides the human in the operator”

Every record is generated by some access procedure. Why pretend that writing $R_A = \mathcal{O}_A(W; \Pi_A)$ escapes interpretation? The operator, selection policy, and reliability model are themselves human constructions.

Reply. It does not escape interpretation; it disciplines it. The point of the notation is precisely that a record is not licensed to identify itself with the world, and a source cue is not licensed to identify itself with reliability. Human construction of an operator does not imply human construction of every constraint on the operator’s success. A bad thermometer is human-made; the temperature-dependent expansion it fails to track is not thereby authored by the thermometer maker. The formalism therefore weakens, rather than strengthens, anthropocentric privilege: humans become one bounded route among others, accountable to what their route preserves and distorts.

11.14. Objection 14: “This is just realism wearing a leather jacket”

The rhetoric is louder than the metaphysics. Of course a realist thinks the world does not wait for us.

Reply. Exactly. The novelty is not the realism. The target is the inconsistency between accepting a mind-independent world in metaphysics and quietly re-centering the human source in epistemic legitimacy. The leather jacket is optional. The contradiction is not.

We removed humanity from the center of the cosmos.
We removed humanity from the center of life.

Epistemology should explain why it deserves the last throne.

12. What is actually being destroyed

Because the rhetoric is intentionally hostile to one familiar picture, the boundaries should be boring.

1. The paper does not claim that AI systems are reliable in general.
2. It does not claim that current AI systems know, believe, understand, testify, or possess moral standing.
3. It does not claim that every fact, trace, or causal relation is literally knowledge under ordinary subject-involving usage.
4. It does claim that the constraint structure to which knowledge must answer is not created merely by the act of knowing it.
5. It does not claim that provenance is normally irrelevant.
6. It does not claim that human accountability, expertise, testimony, or institutions are dispensable.
7. It does not claim that distrust of machines is morally equivalent to testimonial injustice against persons.
8. It does not claim novelty for anti-anthropocentrism as such. Its contribution is the diagnosis of the Possession–Constitution Collapse and its manifestation as knower fetishism in AI-mediated inquiry.

What is being destroyed is narrower and more consequential: the presumption that because a subject must occupy the state of knowing, the recognized knower therefore sits at the metaphysical center of epistemic value.

The subject is a participant in the relation. It arrives into a world already capable of correcting it.

The knower does not confer reality on the evidence.
The evidence confers liability on the knower.

That asymmetry is the whole scandal.

The subject is not the sun around which truth must orbit. At most, the subject is one late-arriving instrument that learned to write grant applications.

13. Conclusion: the knower arrived late

The world was not epistemically blank before someone looked at it. Reality contained differences before anyone named them, lawful dependencies before anyone modeled them, traces before anyone read them, and constraints before anyone argued from them. Before anyone knew, there was already something to misdescribe. If one reserves the word “knowledge” for successful possession by a subject, nothing in this paper objects. Keep the word. The point is that possession arrives late, while answerability does not.

That lateness matters. Once the knower is treated as the owner rather than the creator of epistemic correctness, several confusions become easier to see. A source can matter without being a truth-maker. An expert can deserve deference without possessing intrinsic

epistemic nobility. An institution can earn authority without minting truth. A machine can contribute to a route to knowledge without being a knower. A human can remain accountable without personally generating every epistemic component for which she answers.

Generative AI makes these distinctions difficult to avoid because it removes a convenient historical bundle. The human author used to package generation, intention, endorsement, responsibility, reputation, and credibility in one visible body. Once generation moves elsewhere, the package opens. We are forced to ask which component we actually valued.

Perhaps that is why the disclosure “written by AI” feels philosophically larger than it should. The machine did not merely enter an epistemic practice. It exposed how much of the practice had been organized around an image of the knower.

Written by AI. Still true.

The title does not say that you should believe it.
It asks why authorship made you forget what kind of
question truth is.

A final provocation follows.

Perhaps knowledge, in the possessive sense, needs a
knower.

Truth did not. Constraints did not. Traces did not.
Dependencies did not.

**The knower arrived late, named the building, and
took credit for the architecture.**

AI may not be a knower. Fine.

The fossil was not a knower either. It could still ruin your paper. The detector does not understand your theory. It can still falsify your prediction. The archive has no beliefs. It can still reveal that you misquoted the record.

Reality has veto power without voting rights.

That does not put the human back at the center.

Acknowledgements

The author is grateful to the philosophical traditions on objective knowledge, information, testimony, reliabilism, scientific instruments, distributed cognition, higher-order evidence, and social epistemology that make the argument possible. Responsibility for all claims, formulations, and remaining errors rests with the author.

Author declarations

Funding. No specific grant from any funding agency in the public, commercial, or not-for-profit sectors was received for this work.

Competing interests. The author declares no competing interests.

AI assistance. Generative AI was used as disclosed on the first page for literature mapping, adversarial testing, language editing, and LaTeX drafting. The author accepts responsibility for the final manuscript.

References

- [1] D. Baird, *Thing Knowledge: A Philosophy of Scientific Instruments*. Berkeley: University of California Press, 2004.
- [2] D. Christensen, “Higher-order evidence,” *Philosophy and Phenomenological Research*, vol. 81, no. 1, pp. 185–215, 2010. doi:10.1111/j.1933-1592.2010.00366.x.
- [3] C. A. J. Coady, *Testimony: A Philosophical Study*. Oxford: Clarendon Press, 1992.
- [4] F. I. Dretske, *Knowledge and the Flow of Information*. Cambridge, MA: MIT Press, 1981.
- [5] E. Fricker, “On the metaphysical and epistemic contrasts between human and AI testimony,” *Inquiry*, published online 29 September 2025. doi:10.1080/0020174X.2025.2553297.
- [6] G. Frege, “The Thought: A Logical Inquiry,” *Mind*, vol. 65, no. 259, pp. 289–311, 1956. doi:10.1093/mind/65.1.289.
- [7] O. Freiman, “Analysis of beliefs acquired from a conversational AI: Instruments-based beliefs, testimony-based beliefs, and technology-based beliefs,” *Episteme*, vol. 21, no. 3, pp. 1031–1047, 2024. doi:10.1017/epi.2023.12.
- [8] R. N. Giere, “Distributed cognition in epistemic cultures,” *Philosophy of Science*, vol. 69, no. 4, pp. 637–644, 2002. doi:10.1086/344627.
- [9] A. I. Goldman, “What is justified belief?” in G. S. Pappas, ed., *Justification and Knowledge: New Studies in Epistemology*. Dordrecht: D. Reidel, 1979, pp. 1–25.
- [10] A. I. Goldman, “Experts: Which ones should you trust?” *Philosophy and Phenomenological Research*, vol. 63, no. 1, pp. 85–110, 2001. doi:10.1111/j.1933-1592.2001.tb00093.x.
- [11] J. Hardwig, “Epistemic dependence,” *The Journal of Philosophy*, vol. 82, no. 7, pp. 335–349, 1985. doi:10.2307/2026523.
- [12] P. Humphreys, *Extending Ourselves: Computational Science, Empiricism, and Scientific Method*. New York: Oxford University Press, 2004.
- [13] P. Humphreys, “The philosophical novelty of computer simulation methods,” *Synthese*, vol. 169, no. 3, pp. 615–626, 2009.
- [14] G. Kahane, “Evolutionary debunking arguments,” *Noûs*, vol. 45, no. 1, pp. 103–125, 2011. doi:10.1111/j.1468-0068.2010.00770.x.
- [15] P. Kitcher, “The division of cognitive labor,” *The Journal of Philosophy*, vol. 87, no. 1, pp. 5–22, 1990. doi:10.2307/2026796.
- [16] J. Lackey, *Learning from Words: Testimony as a Source of Knowledge*. Oxford: Oxford University Press, 2008.
- [17] M. Fricker, *Epistemic Injustice: Power and the Ethics of Knowing*. Oxford: Oxford University Press, 2007.
- [18] K. R. Popper, “Epistemology Without a Knowing Subject,” in B. van Rootselaar and J. F. Staal, eds., *Logic, Methodology and Philosophy of Science III*, Studies in Logic and the Foundations of Mathematics, vol. 52. Amsterdam: North-Holland/Elsevier, 1968, pp. 333–373. doi:10.1016/S0049-237X(08)71204-7.
- [19] H. Reichenbach, *The Direction of Time*. Berkeley: University of California Press, 1956.
- [20] S. Sahebi and P. Formosa, “The AI-mediated communication dilemma: Epistemic trust, social media, and the challenge of generative artificial intelligence,” *Synthese*, vol. 205, art. 128, 2025. doi:10.1007/s11229-025-04963-2.
- [21] T. Williamson, *Knowledge and Its Limits*. Oxford: Oxford University Press, 2000.
- [22] L. T. Zagzebski, *Epistemic Authority: A Theory of Trust, Authority, and Autonomy in Belief*. Oxford: Oxford University Press, 2012.